

Lavish Laller

Backend & Systems Engineer | Java, Spring Boot, Rust

lallerlavish2023@gmail.com | +91-8607243184 | [GitHub](#) | [LinkedIn](#) | [LeetCode](#)

Education

B.Tech, Computer Engineering

J.C. Bose UST, YMCA Faridabad, India

2022 – 2026

CGPA: 8.21 / 10

Class XII

Govt. Sr. Sec. School, Ladwa (KKR), India

2022

95.4%

Work Experience

Samsung R&D Institute (SRI-D)

Noida, India

Software Engineering Intern

Jan 2026 – July 2026

- Built an on-device AI daemon in **Rust** with an autonomous, fault-tolerant self-diagnosis pipeline for edge system recovery.
- Automated resolution of **40%** of observed system failures through deterministic recovery workflows.
- Cut p95 initialization latency from **800 ms to 480 ms** by benchmarking and tuning Microsoft **Semantic Kernel** under high-load edge workloads.
- Engineered integrations for **MCP (Model Context Protocol)** servers, enabling zero-install execution of third-party services within the native UX.
- Developed a concurrent execution engine mapping natural language to terminal commands over **gRPC** and **Unix Domain Sockets**.
- Tech Stack:** Rust, LLMs, Semantic Kernel, MCP, gRPC, Unix Domain Sockets, Tizen OS

Projects

mcp-wall / OS-Level Sandbox Runtime for Untrusted MCP Servers

[GitHub](#)

- Kernel-enforced sandbox for untrusted MCP servers using Linux **Landlock** and **seccomp**, blocking the filesystem-escape class of **CVE-2025-53109/53110**.
- 2.4 MB** static binary, **3.6 MB** RAM, **~8 μs** per tool call, **0.5 ms** startup; adds a TOCTOU-safe egress filter and DLP redaction of AWS keys and PII.
- Deterministic adversarial suite of **14 attack scenarios** in CI.
- Tech Stack:** Rust, MCP, Linux Kernel (Landlock, seccomp), JSON-RPC, CI/CD

PaperMind / Multi-Tenant RAG Document Q&A Backend

[GitHub](#)

- Multi-tenant **RAG** service on **Spring AI + pgvector (HNSW, cosine)** answering from user documents with page-level citations; per-user isolation enforced as a **userId** filter on every search and asserted by test.
- 2.0–2.5 s** end-to-end answers streamed over **SSE** with citations in the first frame; **7/8** graded correct across **3** chunking configs on a reproducible retrieval harness.
- Async ingestion off the request thread; **SHA-256** dedupe, type allowlist and **2000**-chunk cap reject documents before embedding spend.
- Tech Stack:** Java 21, Spring Boot, Spring AI, PostgreSQL, pgvector, JWT, Docker, REST, SSE

AlgoForge / Multi-Agent AI System for Human-like Problem Solving

[GitHub](#)

- Multi-agent orchestration engine on Microsoft **AutoGen** with **Gemini** and **OpenAI**, decomposing DSA problems across specialized agent roles.
- 3×** faster end-to-end solve time than a single-agent baseline via parallel role execution.
- Async **FastAPI** backend for concurrent task execution and real-time state, exposed as a REST API with **Firebase** auth.
- Tech Stack:** Python, FastAPI, AutoGen, Gemini API, OpenAI SDK, Firebase, Docker

Technical Skills

Languages: Rust, Python, Java, SQL

Systems: Systems Programming, Linux Kernel Internals (Landlock, seccomp), Process Sandboxing, OS Security

Backend: Spring Boot, Spring AI, FastAPI, REST, gRPC, JSON-RPC, SSE, JWT Auth, PostgreSQL, JPA/Hibernate, Unix Domain Sockets, Tokio, Docker, CI/CD, Git

AI & Agents: RAG, Vector Search (pgvector, HNSW), Embeddings, LLM Agents, Multi-Agent Systems, Model Context Protocol (MCP), Semantic Kernel, AutoGen, On-Device Inference

Fundamentals: Data Structures & Algorithms, Operating Systems, System Design, Concurrency

Honors & Achievements

- Presented on-device AI daemon work to senior leadership visiting from **Samsung headquarters** at an internal AI workshop.
- Contributed data and analysis to a company **town hall** on upcoming technologies and project roadmaps, working with a cross-functional team.
- Solved **500+** problems across LeetCode, Codeforces, and CodeChef.